

Chapter 3

Properties of coefficient estimates

In chapter 1 we described properties of the coefficient estimates, $\hat{\beta}$, in the Gaussian linear model

$$\mathcal{Y} \sim \mathcal{N}(\mathbf{X}\beta, \sigma^2 \mathbf{I}_n).$$

The estimates are called the *least squares estimates* because they minimize the sum of squared residuals from the observed responses, \mathbf{y} . That is,

$$\hat{\beta} = \arg \min_{\beta} S(\beta) = \arg \min_{\beta} \|\mathbf{y} - \mathbf{X}\beta\|^2$$

3.1 Geometric Properties

Recall that $\text{col}(\mathbf{X})$, the *column span* of the $n \times p$ model matrix \mathbf{X} is a linear subspace of the response space, \mathbb{R}^n ,

$$\text{col}(\mathbf{X}) = \{\mathbf{X}\beta : \beta \in \mathbb{R}^p\}$$

The dimension of $\text{col}(\mathbf{X})$ is $k = \text{rank}(\mathbf{X})$ and the QR decomposition used in R uses column pivoting to ensure that the first k columns of \mathbf{Q} are an orthonormal basis for $\text{col}(\mathbf{X})$.

At the risk of some confusion, we will refer to these k columns as \mathbf{Q}_1 which is equivalent to our previous definition in the most common case of full column rank for \mathbf{X} .

The fitted values, $\hat{\mathbf{y}}$, are the (orthogonal) projection of \mathbf{y} onto $\text{col}(\mathbf{X})$

$$\hat{\mathbf{y}} = \mathbf{H}\mathbf{y} = \mathbf{Q}_1\mathbf{Q}_1'\mathbf{y}$$

where the “hat matrix”, $\mathbf{H} = \mathbf{Q}_1\mathbf{Q}_1'$, is a projection matrix of rank

$$\text{tr}(\mathbf{Q}_1\mathbf{Q}_1') = \text{tr}(\mathbf{Q}_1'\mathbf{Q}_1) = \text{tr}(\mathbf{I}_k) = k$$

The diagonal matrix, \mathbf{D} , in the singular value decomposition (sect. 1.4.3, p. 17), $\mathbf{X} = \mathbf{U}_1\mathbf{D}\mathbf{V}'$, has exactly $p - k$ values that are (effectively) zero and these will be in the last $p - k$ positions. (Recall that the singular values, which must be non-negative, are in decreasing order.) Thus the first k columns of \mathbf{U}_1 also form an orthonormal basis for $\text{col}(\mathbf{X})$.

The residual at the parameter estimates, $\hat{\mathbf{e}} = \mathbf{y} - \hat{\mathbf{y}}$ is orthogonal to $\text{col}(\mathbf{X})$. We can prove this by showing that $\hat{\mathbf{e}}$ is orthogonal to the k columns of \mathbf{Q}_1 which form a basis for $\text{col}(\mathbf{X})$.

$$\mathbf{Q}'_1 \hat{\mathbf{e}} = \mathbf{Q}'_1 (\mathbf{y} - \hat{\mathbf{y}}) = \mathbf{Q}'_1 (\mathbf{I}_n - \mathbf{Q}_1 \mathbf{Q}'_1) \mathbf{y} = \left(\mathbf{Q}'_1 - \underbrace{\mathbf{Q}'_1 \mathbf{Q}_1}_{\mathbf{I}_p} \mathbf{Q}'_1 \right) \mathbf{y} = \mathbf{0}$$

This is also an obvious geometric property that to minimize the distance between a point on a hyperplane and a general point in the response space, $\arg \min_{\beta} \|\mathbf{y} - \mathbf{X}\beta\|^2$, you use orthogonal projection of \mathbf{y} onto $\text{col}(\mathbf{X})$ which implies that the residual is orthogonal to $\text{col}(\mathbf{X})$.

Often this relationship is characterized as the *normal equations*. The residual will be orthogonal to $\text{col}(\mathbf{X})$ if it is orthogonal to all the columns of \mathbf{X} , which is to say

$$\mathbf{X}' (\mathbf{y} - \mathbf{X}\hat{\beta}) = \mathbf{0} \Rightarrow (\mathbf{X}'\mathbf{X})\hat{\beta} = \mathbf{X}'\mathbf{y}$$

3.2 Calculus Approach

The function

$$\begin{aligned} S(\beta) &= \|\mathbf{y} - \mathbf{X}\beta\|^2 \\ &= (\mathbf{y} - \mathbf{X}\beta)' (\mathbf{y} - \mathbf{X}\beta) \\ &= \mathbf{y}'\mathbf{y} - \mathbf{y}'\mathbf{X}\beta - \beta'\mathbf{X}'\mathbf{y} + \beta'\mathbf{X}'\mathbf{X}\beta \\ &= \mathbf{y}'\mathbf{y} - 2\beta'\mathbf{X}'\mathbf{y} + \beta'\mathbf{X}'\mathbf{X}\beta \end{aligned}$$

is a real-valued function of the p -vector, β , ($S : \mathbb{R}^p \rightarrow \mathbb{R}$), with gradient vector

$$\frac{dS}{d\beta} = -2\mathbf{X}'\mathbf{y} + 2\mathbf{X}'\mathbf{X}\beta.$$

Thus a critical point, β_c , at which the gradient is zero, satisfies

$$\mathbf{X}'\mathbf{X}\beta_c = \mathbf{X}'\mathbf{y}.$$

The Hessian matrix of $S(\beta)$,

$$\frac{d^2 S}{d\beta d\beta'} = 2\mathbf{X}'\mathbf{X}$$

is positive semi-definite. If \mathbf{X} is full rank then $\mathbf{X}'\mathbf{X}$ is positive definite and the critical point will be the minimizer of $S(\beta)$.

3.3 Algebraic Properties of $\hat{\beta}$

1. $\hat{\beta}$ satisfies $\mathbf{X}'\mathbf{X}\hat{\beta} = \mathbf{X}'\mathbf{y}$ and minimizes $S(\beta) = \|\mathbf{y} - \mathbf{X}\beta\|^2$
2. If \mathbf{X} is of rank p , then $\hat{\beta}$ is unique, satisfying $\mathbf{X}'\mathbf{X}\hat{\beta} = \mathbf{X}'\mathbf{y}$ or, equivalently, $\mathbf{R}\hat{\beta} = \mathbf{Q}'_1\mathbf{y}$ for an invertible, upper-triangular $p \times p$ matrix \mathbf{R} .

3. If $\text{rank}(\mathbf{X}) < p$, $\mathbf{X}'\mathbf{X}\hat{\boldsymbol{\beta}} = \mathbf{X}'\mathbf{y}$ has multiple solutions for $\hat{\boldsymbol{\beta}}$ but $\hat{\mathbf{y}} = \mathbf{X}\hat{\boldsymbol{\beta}}$ is the same for all such $\hat{\boldsymbol{\beta}}$

Proof. To prove item 3: Suppose that $\hat{\boldsymbol{\beta}}_1$ and $\hat{\boldsymbol{\beta}}_2$ are such that $\mathbf{X}'\mathbf{X}\hat{\boldsymbol{\beta}}_1 = \mathbf{X}'\mathbf{X}\hat{\boldsymbol{\beta}}_2 = \mathbf{X}'\mathbf{y}$. Then

$$\mathbf{X}'(\mathbf{X}\hat{\boldsymbol{\beta}}_1 - \mathbf{X}\hat{\boldsymbol{\beta}}_2) = \mathbf{0}$$

which implies that

$$0 = (\hat{\boldsymbol{\beta}}_1 - \hat{\boldsymbol{\beta}}_2)\mathbf{X}'\mathbf{X}(\hat{\boldsymbol{\beta}}_1 - \hat{\boldsymbol{\beta}}_2) = \|\mathbf{X}\hat{\boldsymbol{\beta}}_1 - \mathbf{X}\hat{\boldsymbol{\beta}}_2\|^2 \Rightarrow \mathbf{X}\hat{\boldsymbol{\beta}}_1 = \mathbf{X}\hat{\boldsymbol{\beta}}_2$$

□

3.4 Rank deficient cases and the Moore-Penrose inverse

In practice a rank-deficient model matrix, \mathbf{X} , is handled by two methods

1. Don't create it in the first place, use $I - 1$ *contrasts* for a factor with I levels.
2. Use the pivoted QR decomposition that retains the original order of the columns except that columns whose diagonal elements in \mathbf{R} would be effectively zero are moved to trailing positions.

After that the calculation proceeds as in the full-rank case except that only the first $k = \text{rank}(\mathbf{X})$ columns are used in \mathbf{Q}_1 and \mathbf{R} is taken as the $k \times k$ upper-left submatrix of the calculated $p \times p$ \mathbf{R} .

When discussing the singular value decomposition in Chap. 1, we mentioned the pseudo-inverse or *generalized inverse* of \mathbf{X} , written \mathbf{X}^- , which formally is called the *Moore-Penrose generalized inverse*. If $\text{rank}(\mathbf{X}) = k < p$ then there are $p - k$ singular values of zero (in practice, very close to zero). The SVD is

$$\mathbf{X} = \mathbf{U}_1 \mathbf{D} \mathbf{V}' = \tilde{\mathbf{U}} \tilde{\mathbf{D}} \mathbf{V}'$$

where $\tilde{\mathbf{U}}$ is the first k columns of \mathbf{U}_1 and $\tilde{\mathbf{D}}$ is the first k rows of \mathbf{D} . \mathbf{D}^- , the Moore-Penrose inverse of \mathbf{D} , is also a diagonal matrix with diagonal elements $1/d_{i,i}$, $i = 1, \dots, k$ and zero thereafter. The Moore-Penrose inverse of $\tilde{\mathbf{D}}$ is the first k columns of \mathbf{D}^- . Finally, the Moore-Penrose generalized inverse of \mathbf{X} is

$$\mathbf{X}^- = \mathbf{V} \mathbf{D}^- \mathbf{U}_1' = \mathbf{V} \tilde{\mathbf{D}}^- \tilde{\mathbf{U}}'$$

This is an interesting theoretical tool but in practice it is not necessary to form the SVD in order to solve rank-deficient least squares problems.

3.4.1 Properties of Generalized Inverses

Let \mathbf{A} be an $n \times p$ matrix and \mathbf{A}^- be its $p \times n$ pseudo-inverse. The conditions that \mathbf{A} and \mathbf{A}^- must satisfy are

1. $\mathbf{A}\mathbf{A}^-\mathbf{A} = \mathbf{A}$ (i.e. $\mathbf{A}\mathbf{A}^-$ maps the columns of \mathbf{A} to themselves.)

2. $\mathbf{A}^- \mathbf{A} \mathbf{A}^- = \mathbf{A}^-$ (i.e. $\mathbf{A}^- \mathbf{A}$ maps the columns of \mathbf{A}^- to themselves.)
3. Both $\mathbf{A} \mathbf{A}^-$ and $\mathbf{A}^- \mathbf{A}$ are symmetric

(It is easy to verify these conditions for our case of $\mathbf{X}^- = \mathbf{V} \mathbf{D}^- \mathbf{U}$ where \mathbf{X} is $n \times p$ with $\text{rank}(\mathbf{X}) \leq p \leq n$. In fact, you will do so on a homework assignment.)

Let $\mathbf{H} = \mathbf{A}^- \mathbf{A}$ be the associated projection in \mathbb{R}^p . Then the condition $\mathbf{A} \mathbf{A}^- \mathbf{A} = \mathbf{A}$ implies

1. \mathbf{H} is idempotent because $\mathbf{H} \mathbf{H} = \mathbf{A}^- \mathbf{A} \mathbf{A}^- \mathbf{A} = \mathbf{A}^- \mathbf{A} = \mathbf{H}$.
2. $\mathbf{A} \mathbf{H} = \mathbf{A}$ (just plug in the definition of \mathbf{H}) so $\text{rank}(\mathbf{A}) \leq \text{rank}(\mathbf{H})$. However, we also have $\text{rank}(\mathbf{H}) \leq \text{rank}(\mathbf{A})$ because $\mathbf{H} = \mathbf{A}^- \mathbf{A}$. Thus $\text{rank}(\mathbf{A}) = \text{rank}(\mathbf{H}) = \text{tr}(\mathbf{H})$
3. A general solution of $\mathbf{A} \mathbf{x} = \mathbf{0}$ is

$$\mathbf{x} = (\mathbf{H} - \mathbf{I}_p) \mathbf{z}$$

where \mathbf{z} is any vector in \mathbb{R}^p

$$\begin{aligned} \mathbf{A} \mathbf{x} &= \mathbf{A}(\mathbf{H} - \mathbf{I}_p) \mathbf{z} \\ &= (\mathbf{A} \mathbf{H} - \mathbf{A}) \mathbf{z} \\ &= (\mathbf{A} - \mathbf{A}) \mathbf{z} = \mathbf{0} \end{aligned}$$

4. A general solution to $\mathbf{A} \mathbf{x} = \mathbf{y}$ is

$$\begin{aligned} \mathbf{x} &= \mathbf{A}^- \mathbf{y} + (\mathbf{H} - \mathbf{I}_p) \mathbf{z} \\ \mathbf{A} \mathbf{A}^- \mathbf{A} \mathbf{x} &= \mathbf{A} \mathbf{x} \Rightarrow \mathbf{A} \mathbf{A}^- \mathbf{y} = \mathbf{y} \\ \mathbf{x} &= \mathbf{A} \mathbf{A}^- \mathbf{y} + \underbrace{\mathbf{A}(\mathbf{H} - \mathbf{I}_p) \mathbf{z}}_0 \text{ and } \mathbf{A} \mathbf{A}^- \mathbf{y} = \mathbf{y} \Rightarrow \mathbf{y} = \mathbf{A} \mathbf{x} \end{aligned}$$

$\hat{\boldsymbol{\beta}} = (\mathbf{X}' \mathbf{X})^- \mathbf{X}' \mathbf{y}$ is a particular least squares solution in the rank deficient case. The general solution is

$$\hat{\boldsymbol{\beta}} = (\mathbf{X}' \mathbf{X})^- \mathbf{X}' \mathbf{y} + (\mathbf{H} - \mathbf{I}_p) \mathbf{z}, \quad \mathbf{z} \in \mathbb{R}^p$$

where $\mathbf{H} = (\mathbf{X}' \mathbf{X})^- (\mathbf{X}' \mathbf{X})$.

3.5 Properties of $\hat{\boldsymbol{\beta}}$

In the full-rank Gaussian linear model

$$\begin{aligned} E[\hat{\boldsymbol{\beta}}] &= E[(\mathbf{X}' \mathbf{X})^{-1} \mathbf{X}' \mathbf{y}] \\ &= (\mathbf{X}' \mathbf{X})^{-1} \mathbf{X}' E[\mathbf{y}] \\ &= (\mathbf{X}' \mathbf{X})^{-1} \mathbf{X}' \mathbf{X} \boldsymbol{\beta} \\ &= \boldsymbol{\beta} \end{aligned}$$

which is to say that the least squares estimator is an *unbiased estimator* of β . Furthermore

$$\begin{aligned}\text{Var}(\hat{\beta}) &= \text{Var}((\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathcal{Y}) \\ &= (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\text{Var}(\mathcal{Y})\mathbf{X}(\mathbf{X}'\mathbf{X})^{-1} \\ &\quad \sigma^2(\mathbf{X}'\mathbf{X})^{-1}\end{aligned}$$

R Exercises: Consider the models fit in Chap. 1

```
> lm1 <- lm(optden ~ 1 + carb, Formaldehyde)
> set.seed(1234) # allow for reproducible "random" numbers
> badDat <- within(data.frame(x1=1:20, x2=rnorm(20,mean=6,sd=0.2),
+                             x4=rexp(20,rate=0.02),
+                             y=runif(20,min=18,max=24)),
+                             x3 <- x1 + 2*x2) # create linear combination
> lm2 <- lm(y ~ x1 + x2 + x3 + x4, badDat)
> lm3 <- lm(count ~ spray, InsectSprays)
> lmlst <- list(lm1=lm1, lm2=lm2, lm3=lm3)
> mmlst <- lapply(lmlst, model.matrix)
```

We know that models `lm1` and `lm3` are full-rank but model `lm2` is rank-deficient.

```
> sapply(lmlst, function(fm) c(rank=fm[["rank"]], p=length(coef(fm))))
```

```
      lm1 lm2 lm3
rank    2   4   6
p       2   5   6
```

which is reflected in the diagonal elements of the \mathbf{R} matrices and in the singular values of the model matrices and in the condition number

```
> lapply(lmlst, function(fm) diag(fm[["qr"]][["qr"]]))
```

```
$lm1
[1] -2.4494897  0.6390097
$lm2
[1] -4.472136e+00  2.578759e+01 -8.668932e-01  2.327514e+02  5.179752e-15
```

```
$lm3
[1] -8.485281  3.162278  3.098387  3.000000  2.828427  2.449490
```

```
> lapply(mmlst, function(mm) svd(mm, nu=0, nv=0)[["d"]])
```

```
$lm1
[1] 2.773349 0.564389
$lm2
[1] 3.197628e+02 8.828492e+01 1.396147e+01 1.446151e-01 3.340320e-15
```

```
$lm3
[1] 9.069136 3.464102 3.464102 3.464102 3.464102 1.323169
```

```
> sapply(lmlst, kappa, exact=TRUE)
```

```
      lm1      lm2      lm3
4.913897e+00 1.512149e+17 6.854102e+00
```

For the full-rank models, `lm1` and `lm3`, the pseudo-inverse, \mathbf{X}^- is simply the matrix that creates the estimated coefficients, $\hat{\boldsymbol{\beta}}$ from the observed response vector \mathbf{y} . We can write it in various forms as

$$\mathbf{X}^- = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}' = \mathbf{R}^{-1}\mathbf{Q}'_1 = \mathbf{V}\mathbf{D}^{-1}\mathbf{U}'_1$$

For full-rank models like these the pseudo-inverse, \mathbf{X}^- is unique and

$$\mathbf{X}^-\mathbf{X} = \mathbf{R}^{-1}\underbrace{\mathbf{Q}'_1\mathbf{Q}_1}_{\mathbf{I}_p}\mathbf{R} = \mathbf{I}_p$$

We can verify the conditions for the Moore-Penrose inverse symbolically. For example,

$$\mathbf{X}\mathbf{X}^-\mathbf{X} = \mathbf{Q}_1\underbrace{\mathbf{R}\mathbf{R}^{-1}}_{\mathbf{I}_p}\underbrace{\mathbf{Q}'_1\mathbf{Q}_1}_{\mathbf{I}_p}\mathbf{R} = \mathbf{Q}_1\mathbf{R} = \mathbf{X}$$

or numerically

```
> X <- mmlst[[3]]
> lm3qr <- lmlst[[3]]$qr
> Q1 <- qr.Q(lm3qr)
> R <- qr.R(lm3qr)
> Xpinv <- backsolve(R, t(Q1))
> zapsmall(Xpinv %*% X)
```

```
      (Intercept) sprayB sprayC sprayD sprayE sprayF
[1,]           1         0         0         0         0         0
[2,]           0         1         0         0         0         0
[3,]           0         0         1         0         0         0
[4,]           0         0         0         1         0         0
[5,]           0         0         0         0         1         0
[6,]           0         0         0         0         0         1
```

```
> all.equal(X %*% Xpinv %*% X, X, check.attr=FALSE)
```

```
[1] TRUE
```

```
> all.equal(Xpinv %*% X %*% Xpinv, Xpinv, check.attr=FALSE)
```

```
[1] TRUE
```

For the rank-deficient model, `lm2`, there are many pseudo-inverses.

```

> X <- mmlst[[2]]
> lm2qr <- lmlst[[2]]$qr
> SVD <- svd(X)
> (rr <- lm2qr$rank)                # rank

[1] 4

> (rrind <- seq_len(rr))           # safer than 1:rr

[1] 1 2 3 4

> (dpinv <- c(1/SVD$d[rrind], rep(0, ncol(X) - rr)))

[1] 0.003127319 0.011326963 0.071625693 6.914905158 0.000000000

> str(Xpinv1 <- SVD$v %*% (dpinv * t(SVD$u)))

num [1:5, 1:20] 1.245927 0.048805 -0.055057 -0.061309 -0.000183 ...

> zapsmall(Xpinv1 %*% X)

      (Intercept)      x1      x2      x3 x4
[1,]           1 0.0000000 0.0000000 0.0000000 0
[2,]           0 0.8333333 -0.3333333 0.1666667 0
[3,]           0 -0.3333333 0.3333333 0.3333333 0
[4,]           0 0.1666667 0.3333333 0.8333333 0
[5,]           0 0.0000000 0.0000000 0.0000000 1

> all.equal(X %*% Xpinv1 %*% X, X, check.attr=FALSE)

[1] TRUE

> all.equal(Xpinv1 %*% X %*% Xpinv1, Xpinv1, check.attr=FALSE)

[1] TRUE

> ## An alternative construction is to reduce the SVD components to the first 4 columns
> str(SVDred <- list(d=SVD$d[rrind], u=SVD$u[,rrind], v=SVD$v[,rrind]))

List of 3
 $ d: num [1:4] 319.763 88.285 13.961 0.145
 $ u: num [1:20, 1:4] 0.0262 0.0559 0.155 0.0498 0.1678 ...
 $ v: num [1:5, 1:4] 0.00901 0.11732 0.05349 0.2243 0.96591 ...

> str(Xpinv2 <- with(SVDred, v %*% (1/d * t(u))))

num [1:5, 1:20] 1.245927 0.048805 -0.055057 -0.061309 -0.000183 ...

> all.equal(Xpinv2, Xpinv1)

```



```
[1] TRUE

> ## Finally, we can use a similar construction on the QR decomposition
> ## taking into account the rearrangement of the columns of X
> Xpiv <- X[, lm2qr$pivot]
> str(Xpinv3 <- rbind(backsolve(qr.R(lm2qr)[rrind, rrind], t(qr.Q(lm2qr)[, rrind])), 0))

num [1:5, 1:20] 1.245927 -0.012504 -0.177675 -0.000183 0 ...

> all.equal(Xpiv %*% Xpinv3 %*% Xpiv, Xpiv, check.attr=FALSE)

[1] TRUE

> all.equal(Xpinv3 %*% Xpiv %*% Xpinv3, Xpinv3, check.attr=FALSE)

[1] TRUE
```

The last two constructions show that the Moore-Penrose pseudo-inverse is a matter of collecting the independent columns at the left hand side of the matrix and the linearly-dependent columns on the right hand side, then truncating the decomposition. In other words, if \mathbf{X} is less than full rank then you just find a set of full-rank columns and proceed as before.

R Exercise: (Simulating linear model fits) The `simulate` functions allow us to simulate a matrix of responses based on a fitted model, then fit all the simulated responses in a single call to `lm`. This is much, much faster than any loop-based approach would be.

The result of `simulate` is a named list of response vectors so we drop the names and convert the list to a matrix.

```
> str(Ymat <- data.matrix(unname(simulate(lm1, 10000))))

num [1:6, 1:10000] 0.0843 0.2584 0.4324 0.5263 0.6142 ...
- attr(*, "dimnames")=List of 2
..$ : chr [1:6] "1" "2" "3" "4" ...
..$ : NULL

> fits <- lm(Ymat ~ carb, Formaldehyde)
> str(coefs <- coef(fits))

num [1:2, 1:10000] -0.00201 0.87284 -0.00922 0.90215 0.00369 ...
- attr(*, "dimnames")=List of 2
..$ : chr [1:2] "(Intercept)" "carb"
..$ : NULL
```

Most of the time we want the coefficients to be a data frame instead with columns corresponding to the coefficient names.

```
> str(coefs <- data.frame(t(coef(fits)), check.names=FALSE))
```

```
'data.frame':      10000 obs. of  2 variables:
 $ (Intercept): num  -0.002005 -0.009224 0.003691 -0.001113 0.000131 ...
 $ carb       : num   0.873 0.902 0.887 0.896 0.892 ...
```

Recall that the “true” coefficients for this model are

```
> printCoefmat(coef(summary(lm1)))

              Estimate Std. Error t value Pr(>|t|)
(Intercept) 0.0050857   0.0078337   0.6492   0.5516
carb         0.8762857   0.0135345  64.7444 3.409e-07
```

For an unbiased estimator the mean of the distribution of the estimator should be the parameter value.

```
> sapply(coefs, mean)
```

```
(Intercept)      carb
0.005078325 0.876313251
```

and the standard deviations should be close to the standard errors

```
> sapply(coefs, sd)
```

```
(Intercept)      carb
0.007943042 0.013647517
```

The correlation of sample of coefficient estimates should be close to the value for the fitted model

```
> summary(lm1, corr=TRUE)$correlation
```

```
              (Intercept)      carb
(Intercept)    1.000000 -0.892664
carb          -0.892664  1.000000
```

```
> cor(coefs)
```

```
              (Intercept)      carb
(Intercept)    1.000000 -0.896498
carb          -0.896498  1.000000
```

If, instead, we wish to consider the variance-covariance matrices, we use

```
> vcov(lm1)
```

```
              (Intercept)      carb
(Intercept) 6.136653e-05 -0.0000946449
carb       -9.464490e-05  0.0001831837
```

```
> var(coefs)
```




Figure 3.1: Empirical density plots of coefficient estimates from data simulated according to the estimated parameters in model `lm1`

	(Intercept)	carb
(Intercept)	6.309192e-05	-0.0000971829
carb	-9.718290e-05	0.0001862547

In Fig. 3.1 we show the empirical density plots for the coefficients separately. Alternatively, we could examine the normal Q-Q plots (Fig. 3.2).

We could also plot contours of the estimated 2-dimensional density (Fig. 3.3). The background of the empirical density contours is like a two-dimensional histogram but using hexagonal shaped bins instead of rectangles.




Figure 3.2: Normal quantile-quantile plots of coefficient estimates from responses simulated according to the estimated parameters in model `lm1`.




Figure 3.3: Normal quantile-quantile plots of coefficient estimates from responses simulated according to the estimated parameters in model `lm1`.